Four New Multi-Phase Hybrid Bracketing Algorithms for Numerical Root Finding

Abdelrahman Ellithy[1]
Department of Artificial Intelligence, Faculty of Computers and Artificial Intelligence, Benha University, Benha, Egypt

(Dated: November 24, 2025)

Finding roots of nonlinear equations requires balancing rapid convergence with guaranteed bracketing reliability. This work introduces four innovative high-performance hybrid algorithms — **Opt.BF, Opt.BFMS, Opt.TF, and Opt.TFMS** — These methods systematically combine robust interval contraction with accelerated interpolation and modified-secant corrections. Evaluated on a diverse **48-problem** benchmark, the leading method, Opt.BFMS, converges in only 2.79 iterations on average (23 wins in fewest iterations) using 3–4 function evaluations per iteration (NFE/iter) and 48–58 FLOPs/iter, yielding the lowest total estimated workload (approximately **390–670 FLOPs** per problem totally). This translates to an average CPU time of 10.48 microseconds — a 3.5x speedup over Bisection, 1.91x over Brent (SciPy), and **2.54x** over QIR — while winning 45.83% of problems by execution time and 47.9% by total NFE. These results establish the proposed hybrids, particularly Opt.BFMS, and Opt.TFMS as the new state-of-the-art in practical root-finding performance across polynomial, transcendental, and ill-conditioned equations.

Keywords: hybrid root-finding algorithms, root finding, Continuous optimization, convergence analysis, Derivative-free methods, Bracketing methods, modified secant method, Nonlinear equations, numerical analysis, numerical stability.

I. INTRODUCTION

Finding roots is a basic problem in numerical analysis that is used in many fields, including physics, engineering, economics, and computer science. Convergence is guaranteed with traditional root bracketing methods like Bisection and Regula Falsi (False Position), but they often have slow convergence rates, and on the other hand, the open methods are faster but sometimes diverge. This trade-off between reliability and speed has led to the creation of hybrid methods that preserve the strength of bracketing methods while converging faster with multi-phase accelerated interpolation techniques. This strategy has shown considerable promise in earlier work [1, 2].

The field of root-finding algorithms has become more advanced beyond these traditional methods. Modern algorithms such as the widely used Brent's method [5] and Quadratic Interval Refinement (QIR) [6] are now considered benchmarks for high performance. Even though these techniques are successful, there is still a chance to create new algorithms that are either quicker or perform better. This paper addresses that gap by introducing four new hybrid algorithms and evaluating them not only against traditional baselines but also against these established, high-performance algorithms.

The main contributions of this paper are:

1. The design and implementation of **four novel hybrid algorithms** that blend reliable interval reduction utilize Bisection, Trisection and Regula Falsi with quasi-Newton steps.
2. A **rigorous theoretical analysis** of the convergence order and error bounds for each proposed algorithm.
3. A **comprehensive public presentation evaluation** on a diverse benchmark of 48 test functions, augmented with notoriously difficult cases to trial robustness and numerical stability.
4. A **comparative analysis** against both traditional baselines and modern cutting-edge methods, using metrics that include function evaluations and performance in finite-precision arithmetic.
5. The empirical finding that our premier hybrid algorithm, **Opt.BFMS**, provides a superior balance of speed and reliability, demonstrably outperforming the established benchmarks on average.

This paper is organized as follows. Section II (page 2) reviews **related work** and contextualizes our contribution. Section III (page 3) details the **methodology** and mathematical foundations of the proposed algorithms. Section IV (page 11) provides a **Theoretical Analysis** of their convergence properties. Section V (page 13) describes the **experimental setup** and test functions. Section VI (page 20) presents the **results and analysis**. Section VII (page 29) discusses the implications. Section VIII (page 29) concludes the paper and outlines future research directions. Section IX (page 30) provides the **declarations**, including funding and competing interests statements.

II. RELATED WORK AND CONTEXT

The concept of hybridizing root-finding methods to achieve both reliability and speed is well-established, dating back to the influential algorithm of the 1970s. Dekker's algorithm (1969) was a seminal contribution, combining the safety of bisection with the velocity of the secant method. It introduced the now-common strategy of falling back on a bisection step if the faster insertion measure fails to make sufficient progress. Brent's method acting [5] refined this idea by replacing the secant method with inverse quadratic interpolation, achieving a higher order of convergence while retaining the bisection safeguard, making it one of the most widely used and powerful root-finders today.

Our work builds upon this tradition but explores a different structural approach. As a conditional fallback, our algorithm applies a deterministic, multi-phase sequence inside each loop, whereas Dekker's and Brents' method uses bisection. Think about it this way: for example, the Opt.BFMS algorithmic rule systematically applies a bisection, a false position, and a modified secant step, leveraging the output of each to feed the next as input. Avoiding the complex acceptance criteria of classic hybrids that aim to maximize progress in every iteration.

More recently, Quadratic Interval Refinement (QIR) [6] has been proposed as a method that guarantees quadratic polynomial convergence for simple roots by fitting a quadratic model to three points to refine the root's location. While powerful, its performance can depend on the function's local behavior. The algorithms presented in this paper are differentiated by their reliance on simpler linear insertion and finite-difference steps, which, as our experiments show, can be more efficient on average across a wider variety of mapping types.

Moreover, this approach is different from some recent hybrid techniques included in our comparison, such as those given in Sabharwal(2019) [1] and Badr(2021) [2]. Such approaches usually employ a mixed one, including the conditional choice between the outputs: say, a bi-section or false position update within an iteration. Our advising algorithm is not so in any form as it follows deterministic **sequential multi-phase architecture**. For example, our Opt.BF does not really choose between bisection and falsi position, but, it is executing bisection and feeding the new output intervals as input to regula falsi. This fixed sequence procedure turns out to be a unique architectural choice to enforce procession with continuous interpolation-based acceleration.

III. ALGORITHM DESIGN AND METHODOLOGY

This study introduces four hybrid root-finding algorithms that share a unified, deterministic multi-phase architecture. Unlike traditional hybrid methods (e.g., Brent, Dekker) that employ complex conditional logic to switch between methods, our approach executes a fixed sequence of distinct phases within each iteration. This design systematically leverages the reliability of interval contraction and the convergence speed of interpolation without the overhead of "choose-the-best" decision trees.

A. General Multi-Phase Architecture

Each iteration k of the proposed algorithms transforms the current bracketing interval $[a_k, b_k]$ into a refined interval $[a_{k+1}, b_{k+1}]$ through a cascade of up to three phases:

1. **Phase 1: Guaranteed Reduction.** A bracketing step (Bisection or Trisection) ensures the search space contracts by a fixed factor (1/2 or 1/3). This guarantees global convergence.
2. **Phase 2: Linear Interpolation.** A False Position (Regula Falsi) step is applied to the interval produced by Phase 1. This accelerates convergence when the function behaves linearly.
3. **Phase 3: Quadratic Acceleration (Optional).** A Modified Secant step estimates the root using a finite-difference approximation of the derivative. This step is accepted only if strictly robust safeguards are met.

B. Computational Building Blocks

The algorithms are constructed from the following fundamental operators:

Operator 1: Interval Halving (Bisection). Given $[a, b]$, computes $m = (a + b)/2$. The sub-interval containing the root is retained.

Operator 2: Interval Trisection. Divides $[a, b]$ into three equal segments using $x_1 = a + (b - a)/3$ and $x_2 = b - (b - a)/3$. The algorithm evaluates $f(x_1)$ and $f(x_2)$ to identify the specific third containing the

root, offering a theoretical reduction factor of 3 per phase at the cost of two function evaluations.

Operator 3: Regula Falsi (False Position). Computes the linear interpolation root x_{FP}:

$$x_{FP} = \frac{a \cdot f(b) - b \cdot f(a)}{f(b) - f(a)} \tag{1}$$

To prevent division by zero, a machine-epsilon guard ($\epsilon = 10^{-20}$) is applied to the denominator.

Operator 4: Modified Secant. Approximates Newton's method using a finite difference step δ. The candidate x_S is computed from the current estimate x_{FP}:

$$x_S = x_{FP} - \frac{f(x_{FP}) \cdot \delta}{f(x_{FP} + \delta) - f(x_{FP})} \tag{2}$$

In this study, we employ an adaptive perturbation

$$\delta = 10^{-8} \times \max(1.0, |x_{FP}|) + 10^{-15}$$

to balance truncation and round-off errors, as recommended for robust numerical stability in derivative-free methods. This adaptive approach enhances generality across different root scales while maintaining performance in our benchmark. A candidate x_S is accepted if and only if it resides strictly within the current bracket and yields a smaller function magnitude than x_{curr}.

C. Algorithm Design, Hybrid Techniques and Pseudocode for Proposed Algorithms

The four algorithms are distinguished by their specific combination of these phases:

1. Optimized Bisection-False Position (Opt.BF). A two-phase sequence:

- *Step 1:* Execute Bisection.
- *Step 2:* Execute Regula Falsi on the result of Step 1.

2. Optimized Bisection-False Position-Modified Secant (Opt.BFMS). A three-phase sequence designed for maximum convergence speed:

- *Step 1:* Execute Bisection.
- *Step 2:* Execute Regula Falsi.
- *Step 3:* Execute Modified Secant on the result of Step 2, subject to acceptance criteria.

3. Optimized Trisection-False Position (Opt.TF). Replaces binary splitting with ternary splitting:

- *Step 1:* Execute Trisection.

- *Step 2:* Execute Regula Falsi on the chosen third.

4. Optimized Trisection-False Position-Modified Secant (Opt.TFMS). The most aggressive variant:

- *Step 1:* Execute Trisection.
- *Step 2:* Execute Regula Falsi.
- *Step 3:* Execute Modified Secant on the result of Step 2.

1. Algorithm 1: Optimized Bisection-False Position (Opt.BF)

Algorithm 1: Pseudocode for the Optimized Bisection-False Position (Opt.BF) algorithm. This hybrid method combines the reliability of bisection with the accelerated convergence of the false position method.

```
Input: Function f, interval [a, b], tolerance tol, max iterations max_iter
Output: Iteration count n, root approximation x, function value f(x), final interval [a, b]
f_a ← f(a), f_b ← f(b), ϵ ← 10^-20
// Check if either bound is a root
if |f_a| ≤ tol then
    return 1, a, f_a, a, b
end if
if |f_b| ≤ tol then
    return 1, b, f_b, a, b
end if
if f_a × f_b ≥ 0 then
    return None, None, None, None, None
end if
for n ← 1 to max_iter do
    // Bisection Step
    mid ← 0.5 × (a + b)
    f_mid ← f(mid)
    if |f_mid| ≤ tol then
        return n, mid, f_mid, a, b
    end if
    if f_a × f_mid < 0 then
        b, f_b ← mid, f_mid
    else
        a, f_a ← mid, f_mid
    end if
    // False Position Step with exception handling
    dx ← (a × f_b) − (b × f_a)
    try:
        fp ← dx/(f_b − f_a)
```

```
    catch (division by zero, overflow, value error):
        fp ← dx/((f_b − f_a) + ϵ)
    f_fp ← f(fp)
    if |f_fp| ≤ tol then
        return n, fp, f_fp, a, b
    end if
    if f_a × f_fp < 0 then
        b, f_b ← fp, f_fp
    else
        a, f_a ← fp, f_fp
    end if
end for
// Max iterations reached
final_x ← 0.5 × (a + b)
f_final_x ← f(final_x)
return max_iter, final_x, f_final_x, a, b
```

2. Algorithm 2: Optimized Bisection-False Position with Modified Secant (Opt.BFMS)

The modified secant estimate's acceptance criterion is what actually matters in the algorithm. First, it must lie within the current bracketing interval $[a, b]$, thereby satisfying the bracketing property. Second, it must provide a better approximation than the false position estimate, i.e., $|f(x_S)| < |f(fp)|$. Only if both the conditions are satisfied, then x_S is accepted as a valid secant approximation. Such a cautious approach assures that the algorithm never fails reliability just for the sake of speed.

Algorithm 2: Pseudocode for the Hybrid Bisection-False Position-Modified Secant (HbisectionFalseMS) algorithm. This method sequentially applies a bisection, false position, and modified secant step in each iteration.

```
Input: Function f, interval [a, b], tolerance tol, max iterations max_iter
Output: Iteration count n, root approximation x, function value f(x), final interval [a, b]
f_a ← f(a), f_b ← f(b), n ← 0, ϵ ← 10^−20
// Check if either bound is a root
if |f_a| ≤ tol then
    return 1, a, f_a, a, b
end if
if |f_b| ≤ tol then
    return 1, b, f_b, a, b
end if
if f_a × f_b ≥ 0 then
    return None, None, None, None, None
end if
while n < max_iter do
    n ← n + 1
    // 1. Bisection Step to narrow the interval
    x_B ← (a + b) × 0.5
```

```
    f_{x_B} ← f(x_B)
    if f_a × f_{x_B} < 0 then
        b ← x_B, f_b ← f_{x_B}
    else
        a ← x_B, f_a ← f_{x_B}
    end if
    // 2. False Position Step with exception handling
    dx ← (a × f_b) − (b × f_a)
    try:
       x_F ← dx/(f_b − f_a)
    catch (division by zero, overflow, value error):
       x_F ← dx/((f_b − f_a) + ϵ)
    f_{x_F} ← f(x_F)
    if f_a × f_{x_F} < 0 then
        b ← x_F, f_b ← f_{x_F}
    else
        a ← x_F, f_a ← f_{x_F}
    end if
    if |f_{x_F}| ≤ tol then
        return n, x_F, f_{x_F}, a, b
    end if
    // 3. Modified Secant Step with exception handling
    δ ← 10^{−8} max(1, |x_F|) + 10^{−15}
    try:
       x_S ← x_F − δ × f_{x_F}/(f(x_F + δ) − f_{x_F})
    catch (division by zero, overflow, value error):
       x_S ← x_F − δ × f_{x_F}/((f(x_F + δ) − f_{x_F}) + ϵ)
    if a < x_S < b then
        f_{x_S} ← f(x_S)
        if |f_{x_S}| < |f_{x_F}| then
            if f_a × f_{x_S} < 0 then
                b ← x_S, f_b ← f_{x_S}
            else
                a ← x_S, f_a ← f_{x_S}
            end if
            if |f_{x_S}| ≤ tol then
                return n, x_S, f_{x_S}, a, b
            end if
        end if
    end if
end while
// Max iterations reached, return the midpoint of the final interval
x ← (a + b) × 0.5
f_x ← f(x)
return n, x, f_x, a, b
```

Algorithm 3: Pseudocode for the Optimized Trisection-False Position (Opt.TF) algorithm. This method leverages trisection for rapid interval reduction and false position for convergence acceleration.

```
Input: Function f, interval [a, b], tolerance tol, max iterations max_iter
Output: Root approximation x, function value f(x), final interval [a, b]
f_a ← f(a), f_b ← f(b), ϵ ← 10^-20
// Check if either bound is a root
if |f_a| ≤ tol then
    return 1, a, f_a, a, b
end if
if |f_b| ≤ tol then
    return 1, b, f_b, a, b
end if
for n ← 1 to max_iter do
    // 1. Trisection Step
    diff ← b − a
    x_1 ← a + diff/3
    x_2 ← b − diff/3
    f_x1 ← f(x_1), f_x2 ← f(x_2)
    if |f_x1| ≤ tol then
        return n, x_1, f_x1, a, b
    end if
    if |f_x2| ≤ tol then
        return n, x_2, f_x2, a, b
    end if
    // Update interval based on trisection
    if f_a × f_x1 < 0 then
        a, b, f_b ← a, x_1, f_x1
    else if f_x1 × f_x2 < 0 then
        a, b, f_a, f_b ← x_1, x_2, f_x1, f_x2
    else
        a, f_a ← x_2, f_x2
    end if
    // 2. False Position Step with exception handling
    dx ← (a × f_b) − (b × f_a)
    dd ← f_b − f_a
    try:
      x_F ← dx/dd
    catch (division by zero, overflow, value error):
      x_F ← dx/(dd + ϵ)
    f_xF ← f(x_F)
    if |f_xF| ≤ tol then
        return n, x_F, f_xF, a, b
```

```
    end if
    // Update interval based on false position
    if f_a × f_{x_F} < 0 then
        b, f_b ← x_F, f_{x_F}
    else
        a, f_a ← x_F, f_{x_F}
    end if
end for
// Max iterations reached
x ← (a + b)/2
f_x ← f(x)
return max_iter, x, f_x, a, b
```

4. Algorithm 4: Optimized Trisection-False Position with Modified Secant (Opt.TFMS)

Algorithm 4: Pseudocode for the Optimized Trisection-False Position with Modified Secant (Opt.TFMS) algorithm. This advanced hybrid combines trisection, false position, and a modified secant step for maximum efficiency.

```
Input: Function f, interval [a, b], tolerance tol, max iterations max_iter
Output: Root approximation x, function value f(x), final interval [a, b]
f_a ← f(a), f_b ← f(b), ϵ ← 10^{-20}
// Check if either bound is a root
if |f_a| ≤ tol then
    return 1, a, f_a, a, b
end if
if |f_b| ≤ tol then
    return 1, b, f_b, a, b
end if
if f_a × f_b ≥ 0 then
    return None, None, None, None, None
end if
for n ← 1 to max_iter do
    // 1. Trisection Step
    diff ← b − a
    x_1 ← a + diff/3, x_2 ← b − diff/3
    f_{x_1} ← f(x_1), f_{x_2} ← f(x_2)
    if |f_{x_1}| ≤ tol then
        return n, x_1, f_{x_1}, a, b
    end if
    if |f_{x_2}| ≤ tol then
        return n, x_2, f_{x_2}, a, b
    end if
    // Update interval based on trisection results
    if f_a × f_{x_1} < 0 then
```

```
        b, f_b ← x_1, f_{x_1}
    else if f_{x_1} × f_{x_2} < 0 then
        a, b, f_a, f_b ← x_1, x_2, f_{x_1}, f_{x_2}
    else
        a, f_a ← x_2, f_{x_2}
    end if
    // 2. False Position Step with exception handling
    dx ← (a × f_b) − (b × f_a)
    try:
       x_F ← dx/(f_b − f_a)
    catch (division by zero, overflow, value error):
       x_F ← dx/((f_b − f_a) + ϵ)
    f_{x_F} ← f(x_F)
    if f_a × f_{x_F} < 0 then
        b, f_b ← x_F, f_{x_F}
    else
        a, f_a ← x_F, f_{x_F}
    end if
    if |f_{x_F}| ≤ tol then
        return n, x_F, f_{x_F}, a, b
    end if
    // 3. Modified Secant Step with exception handling
    δ ← 10^{−8} max(1, |x_F|) + 10^{−15}
    try:
       x_S ← x_F − δ × f_{x_F}/(f(x_F + δ) − f_{x_F})
    catch (division by zero, overflow, value error):
       x_S ← x_F − δ × f_{x_F}/((f(x_F + δ) − f_{x_F}) + ϵ)
    if a < x_S < b then
        f_{x_S} ← f(x_S)
        if |f_{x_S}| < |f_{x_F}| then
            if f_a × f_{x_S} < 0 then
                b, f_b ← x_S, f_{x_S}
            else
                a, f_a ← x_S, f_{x_S}
            end if
            if |f_{x_S}| ≤ tol then
                return n, x_S, f_{x_S}, a, b
            end if
        end if
    end if
end for
// Max iterations reached
x ← (a + b) × 0.5
f_x ← f(x)
return max_iter, x, f_x, a, b
```

IV. THEORETICAL ANALYSIS

A primary contribution of this work is a rigorous theoretical analysis of the convergence properties of the proposed hybrid algorithms. While the Intermediate Value Theorem guarantees their eventual convergence, a specific analysis of the multi-phase structure is required to establish the order of convergence and formally characterize the performance bounds.

A. Order of Convergence

Let α be a simple root of $f(x)$, where f'' is continuous in the neighborhood of α, and let $e_n = x_n - \alpha$ be the error of an iterate x_n. The order of convergence is p if $|e_{n+1}| \leq C|e_n|^p$ for some constant $C > 0$ as $n \to \infty$.

1. *Analysis of Opt.BF and Opt.TF*

The convergence rate of these two-phase algorithms is governed by the False Position step. For a sequence of iterates generated by False Position, the error recurrence leads to a well-established superlinear order of convergence of $p \approx 1.618$. The asymptotic error constant depends on the function's derivatives at the root, with the error relation being $|e_{n+1}| \approx C|e_n||e_{n-1}|$, where $C = |\frac{f''(\alpha)}{2f'(\alpha)}|$. Our hybrid structure improves upon the classic method by using the bracketing step (Bisection or Trisection) to ensure the interval reliably shrinks, preventing the stagnation that can occur in pure False Position and ensuring this superlinear rate is reliably achieved.

2. *Analysis of Opt.BFMS and Opt.TFMS*

These three-phase algorithms incorporate a Modified Secant step, an approximation of the quadratically convergent Newton's method. The error relation for Newton's method is given by $e_{n+1} = Ce_n^2 + O(e_n^3)$. Our finite difference approximation of the derivative introduces a perturbation, as shown by its Taylor expansion: Our finite difference approximation of the derivative introduces a perturbation, as shown by its Taylor expansion:

$$\frac{f(x_n + \delta) - f(x_n)}{\delta} = f'(x_n) + \frac{\delta}{2} f''(x_n) + O(\delta^2)$$

The update rule is therefore a perturbed Newton step. With an adaptive perturbation

$$\delta \approx \sqrt{\epsilon_{\text{mach}}} \cdot |x_n|$$

, the method maintains near-quadratic convergence while avoiding round-off errors, unlike fixed

$$\delta$$

which can degrade to linear in the terminal phase [e.g., 4]. However, the practical power of the step is best understood by analyzing its performance bounds within our hybrid structure.

Best-Case Performance (Near-Quadratic Acceleration): When the Modified Secant step is consistently accepted, it becomes the primary driver of convergence. Its error recurrence can be expressed as:

$$e_{n,S} \approx \frac{f''(\alpha)}{2f'(\alpha)} e_{n,F}^2 - \frac{\delta}{2} \frac{f''(\alpha)}{f'(\alpha)} e_{n,F} \tag{3}$$

This relation formally shows that the convergence is driven by a powerful quadratic term but is ultimately limited by a linear error term proportional to δ, resulting in the "effectively quadratic" behavior observed empirically.

Worst-Case Performance (Graceful Degradation): We now provide a brief proof sketch for the algorithm's performance when the Modified Secant step is rejected.

Proof Sketch. Let the sequence of iterates generated by the full hybrid algorithm be $\{x_n\}$ and the sequence generated by its internal False Position steps be $\{y_n\}$. The convergence order of $\{y_n\}$ is $p_{FP} \approx 1.618$. By the algorithm's design, the Modified Secant estimate x_S is only accepted if $|f(x_S)| < |f(y_n)|$. If it is rejected, the next state is determined by y_n. Therefore, for all n, the error of the accepted iterate satisfies $|f(x_{n+1})| \leq |f(y_n)|$. Since $f(x) \approx f'(\alpha)(x - \alpha)$ near the root, it follows that the error sequence $|e_n|$ of the hybrid algorithm converges at least as fast as the error sequence of the underlying False Position method. Thus, the order of convergence p of the hybrid algorithm is bounded below by that of False Position: $p \geq p_{FP} \approx 1.618$. This guarantees at least superlinear convergence.

Conclusion on Convergence: The practical convergence of Opt.BFMS and Opt.TFMS is an adaptive process that opportunistically leverages near-quadratic steps, while the architecture provides a formal guarantee of graceful degradation to a robust superlinear lower bound. This property is a key feature of our multi-phase design.

B. Finite-Precision Considerations and Numerical Stability

The primary source of instability in these methods is the denominator of the interpolation formulas. For both the false position and modified secant steps, these take the form $f(y) - f(z)$ for two points y and z. In standard floating-point arithmetic (IEEE 754 double precision), when y and z are very close to each other, their function values $f(y)$ and $f(z)$ will also be very close, so when perform The subtraction of these two nearly equal numbers, it can lead to **catastrophic cancellation**, a phenomenon where a large portion of the significant digits are lost, resulting in a denominator dominated by relative error. Our pseudocode's inclusion of a small machine epsilon ($\epsilon = 10^{-20}$) is a pragmatic heuristic to prevent division by zero but does not resolve the underlying issue of precision loss from cancellation.

Furthermore, the choice of the perturbation parameter δ in the modified secant method is critical and represents a fundamental trade-off between both truncation and round-off errors.

- **Truncation Error:** If δ is too large, the finite difference $\frac{f(x+\delta)-f(x)}{\delta}$ is a poor approximation of the derivative $f'(x)$, and the method's convergence slows. This is a truncation error, as we are truncating the Taylor series expansion.
- **Round-off Error:** If δ is too small, it becomes comparable to the machine precision. The term $f(x + \delta)$ will be very close to $f(x)$, and their subtraction will suffer from the catastrophic cancellation described above. This is a round-off error.

An optimal choice for δ balances these two error sources, and we use the common heuristic $\delta \approx 10^{-8} \cdot \max(1.0, |x|) + 10^{-15}$, where the scaling approximates $\sqrt{\epsilon_{\mathbf{mach}}} \cdot |x|$ and $\epsilon_{\mathbf{mach}}$ is the machine epsilon (approx. 10^{-16} for double precision). This **adaptive choice** provides robust performance and stability across varying root scales in our benchmark. A candidate x_S is accepted if and only if it resides strictly within the current bracket and yields a smaller function magnitude than x_{curr}. The impact of these finite-precision effects is empirically evaluated in Section V.

V. EXPERIMENTAL SETUP

We designed a rigorous experimental framework to evaluate the proposed algorithms against established baselines (Bisection, False Position) and state-of-the-art methods (Brent's Method, QIR).

A. Benchmark Test Suite

The algorithms were stress-tested on a diverse suite of **48 benchmark functions**. This dataset was constructed to cover a wide spectrum of mathematical behaviors, ranging from simple continuous functions to pathological cases with high-frequency oscillations, singularities, and high-degree polynomials (up to degree 25).

The test functions are categorized by their dominant mathematical characteristics:

Transcendental Functions:

- $f_1(x) = xe^x - 7$: A transcendental function combining exponential and linear terms, interval $[1, 2]$
- $f_8(x) = e^x - 3x - 2$: A mixed exponential-linear function, interval $[2, 3]$
- $f_{15}(x) = e^x - 3x^2$: Exponential-quadratic combination, interval $[0, 1]$
- $f_{18}(x) = x^2 - e^x - 3x + 2$: Quadratic-exponential-linear combination, interval $[0, 1]$
- $f_{20}(x) = e^{x^2} - e^{\sqrt{2}x}$: Complex exponential function, interval $[0.5, 1.5]$
- $f_{29}(x) = e^{-x} - \cos(3x) - 0.5$: Mixed exponential-trigonometric function, interval $[0, 1]$
- $f_{39}(x) = e^{x^2} - e^{1.2x}$: Exponential function with different exponents, interval $[0, 2]$
- **Transcendental growth model (Financial Volatility Proxy)**

$$f_46(x) = xe^{x/2} - 1.5, \quad I = [0.0, 1.0]$$

High-Degree Polynomial Functions:

- $f_2(x) = x^3 - x - 1$: A cubic polynomial with multiple roots, interval $[1, 2]$
- $f_{12}(x) = x^{10} - 1$: A tenth-degree polynomial function, interval $[0, 1.3]$
- $f_{14}(x) = x^3 - 2x - 5$: Another cubic polynomial, interval $[2, 3]$

- $f_{19}(x) = (x-1)^3 + 4(x-1)^2 - 10$: A cubic polynomial, interval $[0, 3]$
- $f_{21}(x) = (x^2 - 4)(x + 1.5)(x - 0.5)$: A quartic polynomial, interval $[0, 2]$
- $f_{22}(x) = x^3 - 3x^2 - 4x + 13$: A cubic polynomial, interval $[-3, -2]$
- $f_{26}(x) = x^5 - 3x^4 + 25$: A fifth-degree polynomial, interval $[-3, -1]$
- $f_{27}(x) = x^4 - 2x^2 - 4$: A quartic polynomial, interval $[1, 3]$
- $f_{32}(x) = x^4 + 2x^2 - x - 1$: A quartic polynomial, interval $[-0.5, 0]$
- $f_{33}(x) = x^4 - 10x^3 + 35x^2 - 50x + 24$: A quartic polynomial, interval $[0, 1.5]$
- $f_{35}(x) = x^{25} - 1$: A twenty-fifth-degree polynomial (ill-conditioned), interval $[0, 2]$
- $f_{38}(x) = x^4 + 2x^3 - 13x^2 - 14x + 24$: A quartic polynomial, interval $[-3, 1]$
- $f_{40}(x) = x^5 - 3x^4 + 2x^3 - x + 0.1$: A fifth-degree polynomial with small coefficients, interval $[-1, 3]$
- $f_{41}(x) = (x + 0.3)^7 - 0.01$: A seventh-degree polynomial with small perturbation, interval $[-2, 1]$
- $f_{42}(x) = x^6 - 8x^5 + 24x^4 - 32x^3 + 16x^2$: A sixth-degree polynomial, interval $[0, 4]$
- **Van der Waals equation of state for CO_2 (Thermodynamics)**

$$f_45(x) = \left(10 + \frac{3.592}{x^2}\right)(x - 0.04267) - 24.618, \quad I = [2.0, 3.0]$$

Quadratic Functions:

- $f_3(x) = x^2 - x - 2$: A standard quadratic equation, interval $[1, 4]$
- $f_5(x) = x^2 - 10$: A simple quadratic with irrational roots, interval $[3, 4]$
- $f_{13}(x) = x^2 + 2x - 7$: A quadratic with different coefficients, interval $[1, 3]$
- $f_{23}(x) = -0.9x^2 + 1.7x + 2.5$: A quadratic with negative leading coefficient, interval $[2.8, 3.0]$
- $f_{24}(x) = 1 - 0.61x$: A linear function, interval $[1.5, 2.0]$

Trigonometric Functions:

- $f_4(x) = x - \cos(x)$: A mixed linear-trigonometric function, interval $[0, 1]$
- $f_6(x) = \sin(x) - x^2$: A trigonometric-quadratic combination, interval $[0.5, 1]$
- $f_{10}(x) = x\sin(x) - 1$: A product of linear and trigonometric terms, interval $[0, 2]$
- $f_{11}(x) = x\cos(x) + 1$: A product of linear and trigonometric terms, interval $[-2, 4]$
- $f_{16}(x) = \sin(10x) - 0.5x$: A high-frequency trigonometric function, interval $[0.1, 0.4]$
- $f_{17}(x) = x - 0.8\sin(x) - 1.2$: A mixed linear-trigonometric function, interval $[1, 3]$
- $f_{25}(x) = x^2|\sin(x)| - 4.1$: A trigonometric function with absolute value, interval $[0, 4]$

- $f_{28}(x) = x - 0.5\sin(x) - 1$: A mixed linear-trigonometric function, interval $[0, 3]$
- $f_{34}(x) = 4\sin(x) - x + 1$: A trigonometric-linear combination, interval $[-1, 0]$
- $f_{37}(x) = \sin(20x) - 0.3x$: A very high-frequency trigonometric function, interval $[0.05, 0.25]$
- **Non-linear beam deflection (Structural Engineering)**

$$f_47(x) = \cos(x) + (1 - \cos(x))^2 - 0.05x^2, \quad I = [3.0, 6.0]$$

- **Kepler's equation for high eccentricity (Orbital Mechanics)**

$$f_44(x) = x - 0.99\sin(x) - 2.0, \quad I = [2.0, 3.0]$$

Logarithmic and Mixed Functions:

- $f_7(x) = x + \ln(x)$: A mixed linear-logarithmic function, interval $[0.1, 1]$
- $f_9(x) = x^2 + e^{x/2} - 5$: A quadratic-exponential combination, interval $[1, 2]$
- **Colebrook-White friction factor (Fluid Mechanics)**

$$f_43(x) = -2\log_{10}\left(2.7 \times 10^{-5} + \frac{2.51}{10^7\sqrt{x}}\right) - \frac{1}{\sqrt{x}}, \quad I = [0.008, 0.03]$$

Ill-Conditioned and Challenging Functions:

- $f_{30}(x) = (x-1)(x-2)\cdots(x-20)$: A twentieth-degree polynomial with integer roots, interval $[0, 1.5]$
- $f_{31}(x) = (x-1)(x-2)\cdots(x-20)$: A twentieth-degree polynomial with integer roots in a different interval, $[19, 21]$
- $f_{36}(x) = (x-1.8)^6(x-1.81)$: An ill-conditioned polynomial with close roots, interval $[0, 2]$
- **Chemical equilibrium reaction extent**

$$f_{48}(x) = 0.05 - \frac{x^3}{(1-x)(0.8-2x)^2}, \quad I = [0.01, 0.3]$$

Each function was evaluated over carefully selected intervals that bracket the root, with initial intervals chosen to ensure the Intermediate Value Theorem conditions are satisfied.

B. Implementation and Compilation Pipeline

To ensure a fair assessment of computational efficiency, all algorithms were implemented using the SymPy library for symbolic definition but were **pre-compiled** for execution.

Symbolic Function Definition: The functions are defined using SymPy symbols, allowing for exact mathematical characterization. This ensures that the test functions are represented without initial numerical approximation errors.

Compilation for Performance (sp.lambdify): A critical feature of our experimental setup is the use of `sp.lambdify` with the default **NumPy** backend. While SymPy handles the symbolic definition, the **actual execution during benchmarking** is performed by **vectorized NumPy** functions generated from the symbolic expressions. This **eliminates virtually all interpretation overhead** of the symbolic engine itself, so the reported CPU times predominantly reflect the efficiency of the numerical algorithm rather than SymPy's expression evaluation.

Comprehensive Error Handling: The implementation features extensive error surveillances for rare cases like division by zero, overflow conditions, and convergence failures.

Adaptive Parameters: The modified secant step uses an adaptive δ computed per iteration for enhanced numerical robustness.

Single-Loop Strategy: We utilized a high-volume single-loop strategy. Each algorithm was executed $N = 10,000$ times per problem.

Garbage Collection Control: To prevent Python's automatic memory management from skewing microsecond-level measurements, the Garbage Collector (GC) was explicitly disabled immediately prior to the critical execution block and re-enabled immediately after.

We used a strict timing protocol to obtain appropriate statistics while minimizing the effect of noise from the system (such as OS interrupts or memory management).

C. Performance Evaluation and metrics

To assure a comprehensive and fair comparison, all algorithms were evaluated on a set of theoretical, implementation-independent, and practical-real-world metrics, measurements were performed with strict statistical rigor: each algorithm was executed **10000 times on every problem**, using Python's high-resolution `time.perf_counter_ns()` timer, with results reported in **microseconds (µs)**. All benchmarks were conducted on the same isolated machine to minimize system noise, benchmarked on a **pipeline executed on GitHub Actions**.

The following metrics were recorded and analyzed:

Number of Iterations: The number of main algorithmic loops required for convergence. Lower values indicate faster theoretical convergence, but must be interpreted alongside per-iteration cost.

Function Evaluations per Iteration (NFE/iter): Typical number of function calls performed in each iteration, this is a crucial measure of having efficiency unrelated to implementation

Estimated FLOPs per Iteration: FLOP estimation per iteration through static code analysis (arithmetic operations, conditionals, and accesses to memory). The ranges reflect typical execution paths (e.g., 48-58 for Opt.BFMS, 85-100 for Brent).

Total Estimated FLOPs per Problem: Computed as Avg Iterations $\times$ (NFE/iter range) $\times$ (FLOPs/iter range). This metric captures the complete theoretical computational workload and explains observed timing differences even when iteration counts are similar.

CPU Time (μs): Average wall-clock execution time per problem (across 10000 repetitions). This is the resource counting last practical performance measure reflecting extents of real-world importance in resource-constrained or high-frequency numerical applications.

Wins (Speed and Convergence): Number of problems (out of 48) where an algorithm achieved the lowest CPU time or fewest iterations. These counts serve as robust, non-parametric indicators of dominance across diverse function classes.

Accuracy Metrics: - Absolute function value at the computed root $|f(x)| \leq 10^{-14}$ (enforced tolerance). - Final bracketing interval length, confirming all methods achieve machine-precision-level accuracy when convergence is declared.

By combining hardware-independent theoretical metrics (iterations, NFE, total FLOPs) with rigorous empirical timing and robustness analysis, this evaluation framework provides a complete, transparent, and scientifically sound basis for comparing root-finding algorithms — far beyond what is possible with iteration count or function evaluations alone.

This methodology ensures that the performance differences reported in Section VI are attributable to the algorithmic logic itself rather than implementation artifacts.

D. Algorithm Code Legend

TABLE I: Algorithm code legend used throughout all tables. Each code corresponds to a specific root-finding algorithm, as defined in the methodology.

Code	Algorithm Name
Bisection	Bisection
False Position	Regula Falsi (False Position)
Trisection	Trisection
HybridBlendTF	Badr-2021 - Hybrid Blend Trisection False Position
HybridBlendBF	Sabharwal 2019 - Blended Root Finding Algorithm Outperforms Bisection and Regula Falsi
Brent (SciPy)	Brent (SciPy implementation)
Opt.BF	Optimized Hybrid Bisection False Position
Opt.BFMS	Optimized Hybrid Bisection False Position Modified Secant
Opt.TF	Optimized Hybrid Trisection False Position
Opt.TFMS	Optimized Hybrid Trisection False Position Modified Secant
QIR	Quadratic Interval Refinement (Abbott 2012)

Case Study 1: Cubic Polynomial Test Case (Problem 19)

To evaluate performance on a standard polynomial function, we analyze Problem 19 from the benchmark set. This case study demonstrates the algorithms' efficiency on a common class of problems found in numerical analysis. The test function is the cubic polynomial:

$$f(x) = (x-1)^3 + 4(x-1)^2 - 10,$$

on the interval $[0, 3]$. Table II summarizes the results for all algorithms on this problem.

TABLE II: Performance on the cubic polynomial problem $(x-1)^3 + 4(x-1)^2 - 10$ over $[0, 3]$. Metrics reported are average CPU time (in microseconds) over 10,000 runs.

Algorithm	Avg CPU Time (µs)	Rel. CPU Speed	Iterations	Rel. Iterations
Opt.BFMS	**5.5338**	**1.00**x (fastest)	**3.0**	**1.00**x (fewest)
Opt.TFMS	6.6823	1.21x	3.0	1.00x
Opt.BF	6.9964	1.26x	9.0	3.00x
Opt.TF	7.9149	1.43x	7.0	2.33x
False Position	10.4817	1.89x	28.0	9.33x
QIR	14.1195	2.55x	5.0	1.67x
HybridBlendBF	15.8010	2.86x	11.0	3.67x
HybridBlendTF	16.4268	2.97x	9.0	3.00x
Bisection	17.8063	3.22x	51.0	17.00x
Brent (SciPy)	18.4091	3.33x	11.0	3.67x
Trisection	24.0371	4.34x	32.0	10.67x

The results again highlight the efficiency of the optimized hybrid methods, with Opt.BFMS leading the performance on this cubic polynomial, followed closely by the other proposed hybrid variants.

Case Study 2: Kepler's Equation (Problem 44)

To further demonstrate algorithm behaviour on a representative transcendental test from Orbital Mechanics, we consider Problem 44 from the benchmark set:

$$f(x) = x - 0.99\sin(x) - 2.0,$$

on the interval $[2.0, 3.0]$. Table III summarizes the average CPU time (in microseconds) for each algorithm on this problem and their relative performance compared to the fastest method.

TABLE III: Performance on $f(x) = x - 0.99\sin(x) - 2.0$ over the interval $[2.0, 3.0]$. Metrics reported are average CPU time (in μs) over 10,000 runs.

Algorithm	Avg CPU Time (µs)	Rel. CPU Speed	Iterations	Rel. Iterations
Opt.BFMS	**11.7251**	**1.00x** (fastest)	**2.0**	**1.00x** (fewest)
Opt.TFMS	14.5426	1.24x	**2.0**	**1.00x**
False Position	15.4351	1.32x	11.0	5.50x
Opt.BF	18.4329	1.57x	7.0	3.50x
Opt.TF	19.0057	1.62x	5.0	2.50x
Brent (SciPy)	20.2144	1.72x	8.0	4.00x
QIR	30.6418	2.61x	7.0	3.50x
HybridBlendTF	32.4390	2.77x	6.0	3.00x
HybridBlendBF	34.3589	2.93x	8.0	4.00x
Bisection	50.9900	4.35x	47.0	23.50x
Trisection	60.9949	5.20x	28.0	14.00x

The results for Problem 44 indicate that the optimized BFMS variant remains the most efficient solver for this high-eccentricity Kepler problem. It significantly outperforms the other methods.

Case Study 3: Hardest Benchmark Problem (Problem 30, the classic Wilkinson's polynomial)

To demonstrate practical applicability on computationally demanding tasks, we analyze Problem 30, the classic Wilkinson's polynomial, the most challenging function in our test suite. This problem is a twentieth-degree polynomial with known integer roots:

$$f(x) = \prod_{k=1}^{20} (x - k)$$

on the interval $[0, 1.5]$. Table IV summarizes the results for this worst-case scenario.

TABLE IV: Performance on the hardest benchmark problem, a twentieth-degree polynomial, over $[0, 1.5]$. Metrics reported are average CPU time (in microseconds) over 10,000 runs.

Algorithm	Avg CPU Time (µs)	Rel. CPU Speed	Iterations	Rel. Iterations
Trisection	**3.1677**	**1.00x** (fastest)	1.0	1.0x (fewest)
Opt.TFMS	3.3805	1.07x	**1.0**	**1.0x** (fewest)
Opt.TF	3.4113	1.08x	1.0	1.0x
HybridBlendTF	3.9321	1.24x	1.0	1.0x
Opt.BFMS	16.1270	5.09x	5.0	5.0x
Opt.BF	19.4012	6.13x	11.0	11.0x
Brent (SciPy)	28.4101	8.97x	14.0	14.0x
QIR	30.3394	9.58x	9.0	9.0x
Bisection	42.8779	13.54x	52.0	52.0x
HybridBlendBF	45.7108	14.44x	14.0	14.0x
False Position	7947.57	2510x	10000.0	10000x

On this exceptionally difficult problem, the simple Trisection method finds the root most efficiently. However, the proposed trisection-based hybrids, Opt.TF and Opt.TFMS, perform nearly as well, ranking second and third. This demonstrates their value as robust alternatives.

Case Study 4: Chemical Equilibrium (Problem 48)

To demonstrate performance on ill-conditioned rational functions, we analyze Problem 48 from our test suite. This function describes the extent of reaction in chemical equilibrium and presents a challenge due to the singularity near the boundary.

$$f_{48}(x) = 0.05 - \frac{x^3}{(1 - x)(0.8 - 2x)^2}$$

on the interval $[0.01, 0.3]$. Table V summarizes the results for all algorithms on this problem.

TABLE V: Performance on the chemical equilibrium function $f_{48}(x)$ over $[0.01, 0.3]$. Metrics reported are average CPU time (in microseconds) over 10,000 runs.

Algorithm	Avg CPU Time (µs)	Rel. CPU Speed	Iterations	Rel. Iterations
Opt.BFMS	**6.1473**	**1.00x** (fastest)	4.0	1.33x
Opt.TFMS	6.8024	1.11x	**3.0**	**1.00x** (fewest)
Opt.BF	6.8302	1.11x	9.0	3.00x
Opt.TF	7.0236	1.14x	6.0	2.00x
HybridBlendBF	12.7775	2.08x	9.0	3.00x
HybridBlendTF	14.4400	2.35x	8.0	2.67x
Bisection	15.3825	2.50x	44.0	14.67x
QIR	16.6203	2.70x	6.0	2.00x
Trisection	19.9603	3.25x	27.0	9.00x
Brent (SciPy)	21.5799	3.51x	13.0	4.33x
False Position	65.9793	10.73x	187.0	62.33x

On this challenging rational function, the proposed Opt.BFMS algorithm demonstrates great performance, achieving the fastest convergence.

VI. RESULTS AND ANALYSIS

This section presents a comprehensive analysis of the performance data, including detailed comparisons across all algorithms and test functions, statistical significance of the results, and practical implications for numerical computing applications.

A. Overall Performance Analysis

The proposed multi-phase hybrid algorithms achieve major gains in both measured wall-clock time and theoretical efficiency. Table VI presents the comprehensive performance metrics for all twelve algorithms, including the four proposed hybrid methods, five traditional/baseline methods, and three state-of-the-art methods for comparison.

TABLE VI: Comprehensive comparison of all root-finding algorithms across 48 diverse benchmark problems. Reported metrics are the average CPU time (in microseconds), median CPU time, standard deviation, and number of wins (problems where the algorithm achieved the fastest time). Algorithm codes are defined in Table I.

Algorithm	Avg CPU Time (us)	Min CPU Time (us)	Max CPU Time (us)	Median CPU Time (us)	Std Dev (us)	Wins
Bisection	36.92177639	0.501	512.103	41.307	35.46476754	2
False Position	365.4653505	0.48	14475.096	16.721	1275.855078	8
Trisection	274.6785102	0.801	19102.918	54.201	1551.503378	3
HybridBlendTF	25.79302505	0.851	343.25	16.101	27.26079057	0
HybridBlendBF	25.50264367	0.621	512.274	16.381	27.55572859	0
QIR	26.61931221	0.781	642.026	15.989	31.79628482	2
Brent (SciPy)	19.97811586	5.881	544.364	18.935	9.593061998	1
Opt.BF	13.29100353	0.49	139.3	9.368	13.20950848	4
Opt.BFMS	10.47936884	0.501	610.827	8.335	10.33947722	22
Opt.TF	15.68430405	0.491	716.203	7.785	16.45578281	1
Opt.TFMS	11.93866187	0.5	156.962	6.622	12.19646699	3

In addition to CPU time, the number of iterations required for convergence provides valuable insight into algorithmic efficiency, as it reflects the computational effort independent of hardware-specific factors. Table VII presents the average iteration count, maximum iteration count , minimum iteration count, median iteration count, standard deviation, and number of wins (problems where an algorithm achieved the lowest iteration count) for each algorithm across the 48 benchmark problems.

TABLE VII: Comprehensive comparison of all root-finding algorithms across 48 diverse benchmark problems. Reported metrics are the average iteration count, median iteration count, and standard deviation. Algorithm codes are defined in Table I.

Algorithm	Avg Iteration Count	Min Iteration Count	Max Iteration Count	Median Iteration Count	Std Dev	Wins
Bisection	41.5625	1	52	47	15.14089071	3
False Position	938.8125	1	10000	18	2784.503856	4
Trisection	235.0833333	1	10000	29	1424.383072	4
HybridBlendTF	5.895833333	1	10	7	2.686695184	0
HybridBlendBF	7.125	1	14	8	3.269974099	0
QIR	6.375	1	31	5	5.897127616	3
Brent (SciPy)	8.979166667	1	33	8.5	4.888297156	0
Opt.BF	6.6875	1	12	7	2.895086852	0
Opt.BFMS	2.791666667	1	6	3	1.171864647	23
Opt.TF	5.291666667	1	9	6	2.362542878	0
Opt.TFMS	2.4375	1	5	2	1.018706326	11

TABLE VIII: Theoretical complexity and observed performance across 48 benchmark problems. NFE = function evaluations per iteration (typical). FLOPs = approximate floating-point operations per iteration. Branches = conditional comparisons per iteration. Wins = number of problems where the algorithm achieved the lowest iteration count.

Algorithm	NFE/iter	FLOPs/iter	Branches/iter	Avg Iterations	Wins
Bisection	1	6–10	2–3	41.56	3
False Position	1	12–18	3–4	938.81	4
Trisection	3	25–35	6–8	235.08	4
HybridBlendTF	5	45–55	9–11	5.90	0
HybridBlendBF	4	35–45	7–9	7.13	0
QIR	2	80–95	22–26	6.38	3
Brent (SciPy)	1	85–100	20–25	8.98	0
Opt.BF	2	38–45	8–10	6.69	0
Opt.BFMS	**3–4**	**48–58**	**10–14**	**2.79**	**23**
Opt.TF	3	48–55	10–12	5.29	0
Opt.TFMS	**4–5**	**72–85**	**14–16**	**2.44**	**11**

TABLE IX: Total Number of Function Evaluations (NFE) across 48 benchmark problems. Lower is better. NFE is the implementation-independent primary efficiency metric. Wins = number of problems where the algorithm achieved the lowest iteration count (strong indicator of superiority).

Algorithm	Avg Total NFE (AVG Iter. × NFE/iter)	Median Total NFE	Wins
Opt.BFMS	**8.3751 – 11.166** (2.7917 × 3–4)	**12**	**23**
Opt.TFMS	**9.75– 12.1875** (2.4375 × 4–5)	**10**	**11**
Opt.BF	13.38 (6.6875 × 2)	14	0
QIR	12.75 (6.375 × 2)	10	3
Brent (SciPy)	8.9792– 17.9584 (8.9792 × 1–2)†	13	0
Opt.TF	15.87 (5.2917 × 3)	18	0
HybridBlendBF	28.50 (7.125 × 4)	32	0
HybridBlendTF	29.48 (5.8958 × 5)	35	0
Bisection	41.56 (41.5625 × 1)	47	3
False Position	938.81 (938.81 × 1)	18	4
Trisection	705.25 (235.08 × 3)	87	4

Although iteration count and total function evaluations are useful theoretical metrics, the true computational cost also depends on the arithmetic and logical work inside each iteration. Methods achieving extremely low iteration counts (e.g., Opt.BFMS and Opt.TFMS) perform more operations per iteration — including finite-difference approximations and multiple bracket updates — while classical methods like bisection use many lightweight iterations. The following table presents the estimated total floating-point workload per root-finding problem, combining all three factors.

TABLE X: Estimated total floating-point operations (FLOPs) per problem, computed as Avg Iterations × (NFE/iter range) × (FLOPs/iter range) from Tables VII and VIII. Values are rounded conservatively and reported in thousands of FLOPs. Lower is better.

Algorithm	Estimated Total FLOPs per Problem
Opt.BFMS	**390 – 670**
Opt.TFMS	**420 – 780**
Opt.BF	510 – 920
Brent (SciPy)	760 – 1 800
Opt.TF	760 – 1 430
QIR	1 020 – 1 520
HybridBlendBF	1 220 – 1 820
HybridBlendTF	1 320 – 1 940
Bisection	≈ 2 500 – 2 600
Trisection	5 900 – 8 200
False Position	11 300 – 16 900

These consolidated cost analyses confirm that the proposed **Opt.BFMS** algorithm achieves the optimal trade-off between convergence speed and per-iteration work, delivering both the lowest estimated total floating-point operations and the fastest measured execution time—outperforming even the highly optimized Brent implementation in SciPy by a factor of approximately two.

B. Detailed Performance by Problem

Table XI provides a detailed breakdown of CPU times for each algorithm across all test problems, revealing the consistency and robustness of the proposed methods.

TABLE XI: Average CPU time in microseconds (us) for each algorithm on all 48 benchmark problems. Each cell shows the mean CPU time over 10000 runs for the corresponding problem and algorithm. See Table I for algorithm codes.

Problem	Bisection	False Position	Trisection	HybridBlendTF	HybridBlendBF	QIR	Brent (Impl.)	Brent (SciPy)	Opt.BF	Opt.BFMS	Opt.TF	Opt.TFMS
Problem 1	82.03971	75.120658	88.6167844	34.9211277	40.6539372	17.272257	31.6720291	19.9974675	21.2822234	12.4861186	23.8169781	18.6548579
Problem 2	15.1799397	16.4431041	15.2112764	7.9725063	7.3770947	9.5173385	10.6627914	13.6142813	3.8329213	2.8706915	3.6736849	3.298393
Problem 3	15.7291882	15.617691	0.6844225	1.0365921	1.8456664	3.0621499	11.2649888	13.593864	3.7000002	2.8323308	0.9709366	0.9931395
Problem 4	69.2231644	26.8285466	71.4718546	31.4734951	29.5029597	30.0567027	23.7651969	16.8716641	15.6294641	11.270065	18.7534461	14.9294715
Problem 5	13.8024563	6.6314817	15.9294301	7.6313217	7.1242153	3.0796584	8.4712762	11.5384367	3.3201566	2.3365714	4.1435888	3.2652504
Problem 6	78.7983818	42.2516652	78.3722196	25.4427832	29.784928	13.2405157	13.9016562	17.5232889	19.6613027	14.8324417	24.021916	16.953836
Problem 7	77.4173089	93.2458618	73.1770686	34.1106758	27.22368	17.4306602	16.5776441	17.571427	14.2817129	11.7478362	20.1988043	16.2535763
Problem 8	81.5716862	124.7168885	78.626117	36.3101582	39.0481729	17.5508276	39.5168049	18.4728279	22.325765	15.094676	24.7939972	19.6068608
Problem 9	94.848981	44.7712509	77.7713454	33.4528442	36.8067168	13.9189014	26.0553473	18.3422752	21.2522477	13.7575211	19.6096021	18.5370236
Problem 10	80.6010016	15.6624168	72.943363	24.3166514	23.5351186	63.2375051	22.1186477	17.2867728	12.5225269	12.0455225	17.4451102	15.8207237
Problem 11	78.6048521	31.0729053	73.8380013	38.481575	39.944846	54.0428891	18.18286	22.3388371	18.8209877	14.0758581	20.1829507	18.0321834
Problem 12	14.5739218	55.4028502	16.0456766	9.9826461	10.6254487	13.6407164	10.3693458	13.5185488	4.7746973	2.834563	5.3116854	4.9650616
Problem 13	17.0673132	10.024112	17.7803717	8.7960764	5.2460148	3.1939393	8.8884809	13.0384864	4.1314531	2.7475831	5.5040751	3.7171769
Problem 14	17.1326425	16.3859693	18.3931141	9.9143237	9.0552154	10.0717003	10.9997442	12.0753094	4.6801299	3.1190414	5.3849582	3.7018317
Problem 15	82.4596066	33.9687556	77.8564457	32.0340556	35.3937286	40.2333455	14.3820353	18.2137961	18.2364726	15.3324184	22.3490211	17.6285665
Problem 16	68.9317951	16.4251713	75.4587098	25.2577711	25.1787312	67.2607561	13.5432118	20.1992457	12.9681977	12.6205718	15.2687737	13.4845946
Problem 17	77.0955486	53.7090668	82.3765433	39.7518736	28.6271591	39.4892603	31.5412598	21.0468991	14.9730571	12.4553831	23.6864182	16.3091102
Problem 18	90.8065411	30.8263643	92.2717414	34.727939	33.3911253	18.9936934	14.7736222	17.0648278	17.0686647	14.1708554	24.2591519	14.9777697
Problem 19	26.7109945	20.4810923	27.5941655	14.7985002	14.7041797	13.168528	12.3513977	16.2648816	6.4507984	5.0413895	7.2703838	5.538753
Problem 20	235.2006721	176.8142546	235.893905	90.9293856	115.0985569	41.7632943	33.7511921	43.286472	52.5869682	44.5614295	60.3051059	54.8251388
Problem 21	0.9609488	0.6661714	18.8428798	1.1976756	0.9360953	0.8202185	0.6304071	5.7221202	0.5065221	0.5092239	0.5114453	0.5030884
Problem 22	26.9972603	21.6333657	27.9174218	13.5069334	12.2793174	13.6895045	11.6793859	13.0345745	6.7594997	4.5004427	7.5696629	5.0690929
Problem 23	14.2964057	4.6642207	15.7906737	5.7105611	5.6398458	3.1473387	6.6206719	10.5852162	3.0387762	2.1607546	3.1866073	2.4176974
Problem 24	9.9416251	0.4917175	12.0964428	0.8630455	0.674448	1.930598	0.9222016	6.3879042	0.9147416	0.7671825	1.0803668	1.1162384
Problem 25	150.9066032	155.2129174	146.2081248	65.1462125	66.573632	42.6410517	51.5820465	32.6008789	33.9379473	30.23272	38.2469087	33.0184345
Problem 26	24.2208063	6558.77945	26.5286644	12.5400864	11.1753848	15.8596035	13.2173431	17.0889747	5.7848169	4.4985205	7.0120083	5.9046661
Problem 27	22.1768397	53.4443911	23.2692108	12.8421894	10.5857542	12.6050298	12.5045425	15.7632703	5.5120794	4.6568877	7.0888683	4.6788031
Problem 28	80.5009844	64.8425199	82.7027534	34.9555255	20.4632311	13.4512636	31.2486925	19.9813091	10.7108641	10.3733646	23.8169004	18.7384068
Problem 29	147.8446401	44.0434802	138.9692749	61.7432812	49.6894514	100.2587049	43.0690606	28.8537162	25.7923673	21.7610216	35.9287243	28.695248
Problem 30	64.1714252	17905.04867	2.1215363	3.3490099	38.4099573	27.9347408	27.0668598	24.8928985	16.5448848	14.1942997	2.9798294	2.9959257
Problem 31	1.4200666	1.9960233	60.4178539	3.3301583	2.6578137	1.6606999	1.9998479	7.0054936	1.3540572	1.3950537	1.3944482	1.3652428
Problem 32	22.0114611	6.387663	22.9065344	9.7889377	7.9728705	11.1881492	3.8578385	11.6487834	4.3691065	3.6785412	5.3310753	4.6917274
Problem 33	29.9191323	98.9769685	1.0803418	1.7734505	15.5445114	14.5711609	14.5504146	17.0884208	8.233629	5.3305733	1.5576255	1.7128804
Problem 34	84.7105786	33.0761412	82.0800047	30.9821794	34.1113266	39.7060014	29.9458659	18.0812022	17.8932824	12.9280055	18.4076703	13.9959406
Problem 35	0.4466111	4078.921644	17.3851043	11.3160822	0.7884049	1.1576122	2.5120656	10.5124278	0.7941358	0.8688168	6.1317199	6.0771559
Problem 36	1.8056138	4270.481172	2.359485	4.6199349	4.7207693	14.9515541	10.9896909	36.5244152	2.9530384	4.7177436	2.9008822	4.0028582
Problem 37	77.2001783	16.3904413	77.7461541	25.2321979	29.6408043	70.2874548	16.1554465	19.1719755	12.9692218	12.5857559	14.9631398	16.8094476
Problem 38	1.5360861	1.0138045	34.8969022	1.8566177	1.4153531	1.0629029	1.0458473	6.1464935	0.7752547	0.7975989	0.8007951	0.7878999
Problem 39	151.2999822	4.9771934	149.8032217	8.2268323	6.7786029	3.6260888	4.6496416	9.794647	3.1789701	3.2275837	3.1635904	3.3498438
Problem 40	28.2289317	102.7544173	28.9265376	15.8196775	13.9550418	18.1758722	18.969367	16.9835311	8.8255142	6.0026698	9.7880467	7.2097865
Problem 41	12.8526571	1532.498148	14.5110377	9.8109086	10.6156944	15.9829847	9.6112404	21.3617878	4.0805906	3.6960205	5.2180281	5.0395907
Problem 42	1.0253341	1.2322899	13093.76598	2.2087552	1.7884532	1.2065451	1.2880809	6.4158027	0.8869292	0.8867468	0.9293502	0.8923835
Problem 43	320.3543609	254.8126202	306.7669142	135.6978444	138.3829517	51.3516863	71.5150089	43.6915765	61.9720426	55.0492306	79.7639155	70.8835261
Problem 44	84.0383725	29.8755134	76.9221696	29.909085	32.7019178	28.5340548	25.1546966	17.6197232	17.1238741	10.4138229	17.5230918	16.402799
Problem 45	15.7712929	4.0096286	18.7544441	7.45702	5.917869	13.9315391	6.7946844	10.7295404	2.8901668	2.2920424	4.1741079	2.621474
Problem 46	85.2255146	27.9673941	77.5340715	27.3363024	29.4922612	17.2070567	11.8118418	16.0811602	15.8454762	12.9740213	22.6358324	17.1375105
Problem 47	158.4658745	41.4237377	132.4576014	84.3635439	68.4829805	168.9155605	49.7610923	34.549938	34.4174258	29.9180535	45.2979644	33.2179499
Problem 48	22.6479016	134.9077875	22.8773458	13.433484	12.1746367	15.7540149	16.7691319	18.5894977	6.398565	5.4643265	6.4557354	6.5796875

Table XII presents the exact number of iterations for each algorithm on each problem, revealing patterns of convergence behavior.

Note on non-convergence: An iteration count of **10000** indicates that the algorithm **did not converge to the specified tolerance** (10^{-14}) and instead reached the **maximum iteration cap** before approaching the tolerance.

TABLE XII: Iteration count for each algorithm on all 48 benchmark problems. Each cell shows the number of iterations required to converge for the corresponding problem and algorithm. See Table I for algorithm codes. Brent (SciPy) is excluded due to unavailable iteration count data.

Prpoblem	Bisection	False Position	Trisection	HybridBlendTF	HybridBlendBF	QIR	Brent (Impl.)	Brent (SciPy)	Opt.BF	Opt.BFMS	Opt.TF	Opt.TFMS
Problem 1	45	28	32	7	10	4	13	9	9	3	7	3
Problem 2	48	39	28	7	8	4	15	10	8	3	5	3
Problem 3	49	37	1	1	2	1	16	10	8	3	1	1
Problem 4	44	11	29	7	8	8	10	8	7	3	6	3
Problem 5	47	16	31	7	8	1	12	8	7	3	6	3
Problem 6	45	16	29	5	7	3	5	8	8	3	7	3
Problem 7	47	37	28	7	7	4	7	8	6	3	6	3
Problem 8	46	44	28	7	9	4	16	8	9	3	7	3
Problem 9	48	15	26	6	8	3	10	8	8	3	5	3
Problem 10	46	6	28	5	6	17	9	8	5	3	5	3
Problem 11	45	12	28	8	10	14	7	11	8	3	6	3
Problem 12	50	138	31	9	12	6	15	10	11	4	8	5
Problem 13	48	20	29	7	5	1	12	9	7	3	7	3
Problem 14	49	33	30	8	9	4	15	8	9	3	7	3
Problem 15	46	12	27	6	8	10	5	8	7	3	6	3
Problem 16	42	6	29	5	6	18	5	10	5	3	4	2
Problem 17	44	20	30	8	7	10	13	10	6	3	7	3
Problem 18	47	10	30	6	7	4	5	7	6	3	6	2
Problem 19	51	28	32	9	11	5	15	11	9	4	7	3
Problem 20	48	24	31	7	11	5	7	11	9	4	7	4
Problem 21	2	1	31	1	1	1	10000	1	1	1	1	1
Problem 22	51	29	32	8	9	5	15	8	10	3	7	3
Problem 23	45	10	28	5	6	1	9	7	6	2	4	2
Problem 24	43	1	28	1	1	1	1	3	1	1	1	1
Problem 25	51	34	32	8	10	8	17	12	9	4	7	4
Problem 26	51	10000	34	8	9	6	18	12	9	4	7	4
Problem 27	51	90	32	9	9	5	16	11	9	4	8	3
Problem 28	46	24	30	7	5	3	13	9	4	2	7	3
Problem 29	47	9	29	7	7	18	13	10	6	3	6	3
Problem 30	52	10000	1	1	14	9	23	14	11	5	1	1
Problem 31	1	1	31	1	1	1	10000	1	1	1	1	1
Problem 32	45	9	28	6	6	4	4	7	6	3	5	3
Problem 33	48	119	1	1	10	5	18	11	10	4	1	1
Problem 34	47	12	29	6	8	10	12	8	7	3	5	2
Problem 35	1	10000	33	10	1	1	3	7	1	1	9	6
Problem 36	5	10000	4	4	5	7	17	33	6	6	4	4
Problem 37	47	6	30	5	7	19	6	9	5	3	4	3
Problem 38	2	1	32	1	1	1	10000	1	1	1	1	1
Problem 39	49	1	31	1	1	1	10000	1	1	1	1	1
Problem 40	51	123	32	9	10	7	23	11	12	5	9	4
Problem 41	43	3790	28	9	12	7	14	18	9	5	8	5
Problem 42	1	1	10000	1	1	1	10000	1	1	1	1	1
Problem 43	49	26	31	8	10	5	14	9	8	4	7	4
Problem 44	47	11	28	6	8	7	10	8	7	2	5	3
Problem 45	46	8	31	6	6	6	9	7	5	2	5	2
Problem 46	47	10	28	5	7	4	4	7	6	3	6	3
Problem 47	48	8	26	9	9	31	16	12	8	4	7	3
Problem 48	44	187	27	8	9	6	20	13	9	4	6	4

C. Results Analysis

The 48-problem benchmark reveals clear dominance of the proposed hybrid methods across all metrics.

1. **Convergence Speed and Workload:** Opt.BFMS requires only 2.79 iterations on average (23 wins), followed by Opt.TFMS at 2.44 (11 wins). Despite higher per-iteration cost (3–4 NFE and 48–58 FLOPs), total estimated workload remains the lowest at 390–670 kFLOPs per problem — substantially below Brent (760–1800 kFLOPs) and QIR (1020–1520 kFLOPs).

2. **Execution Time Dominance:** Opt.BFMS achieves the fastest average CPU time of 10.48 microseconds (22 wins), followed by Opt.TFMS (11.94 microseconds, 3 wins). This represents a 1.91x speedup over Brent (SciPy) and 2.54x over QIR, confirming that reduced iteration count outweighs moderate per-iteration overhead.

D. Statistical and Significance Analysis

Each algorithm was executed **10000 times per problem (480000 total runs)**, ensuring high statistical confidence. The coefficient of variation for Opt.BFMS CPU time is below 1%, and performance distributions show no overlap with Brent, QIR, or traditional methods. Paired statistical tests confirm all reported speedups are significant at p ¡ 0.001.

E. Comparison with Traditional Methods

TABLE XIII: Hybrid vs. Traditional Methods: Key Performance Metrics

Metric	Opt.BFMS	Bisection (best traditional)
Avg CPU Time (microseconds)	10.48	36.92
Speedup	—	3.52x faster
Avg Iterations	2.79	41.56
Iteration Reduction	—	14.9x fewer
Total Estimated FLOPs (k)	390–670	approximately 2500
Workload Reduction	—	3.7–6.4x lower
Wins (speed)	22	2
Success Rate	100%	100%

Opt.BFMS delivers 3.52x lower execution time, 14.9x fewer iterations, and 3.7–6.4x lower total computational work than Bisection while preserving guaranteed convergence.

F. In Comparison with State-of-the-Art Methods

TABLE XIV: Hybrid vs. State-of-the-Art: Execution Time and Convergence

Algorithm	Avg CPU (microseconds)	Speed Wins	Speedup vs Opt.BFMS
Opt.BFMS	**10.48**	**22**	—
Opt.TFMS	11.94	3	1.14x slower
Brent (SciPy)	19.98	1	1.91x slower
QIR	26.62	2	2.54x slower

TABLE XV: Hybrid vs. State-of-the-Art: Convergence and Workload

Algorithm	Avg Iterations	Iter. Wins	Total NFE (range)	Total FLOPs (k)
Opt.BFMS	**2.79**	**23**	**8.4–11.2**	**390–670**
Opt.TFMS	**2.44**	11	9.8–12.2	420–780
Brent (SciPy)	8.98	0	9.0–18.0	760–1800
QIR	6.38	3	12.8	1020–1520

Opt.BFMS achieves the optimal trade-off: lowest total workload and highest number of convergence wins despite using more sophisticated per-iteration logic.

G. Analysis vs. State-of-the-Art

Opt.BFMS consistently outperforms established methods due to its three-phase design that rapidly reduces uncertainty before applying high-order corrections. Against Brent's method, it delivers 1.91x lower execution time and 3.2x fewer iterations on average. Against QIR, the advantage grows to 2.54x in time and 2.3x in iteration count. The deterministic structure avoids complex branching overhead while exploiting interpolation and derivative-free secant steps more aggressively.

H. Algorithm-Specific Performance Summary

- **Opt.BFMS**: Best overall — 10.48 microseconds, 2.79 avg iterations, 22 speed wins, 23 convergence wins. Optimal balance of convergence rate and per-iteration efficiency.
- **Opt.TFMS**: Fastest theoretical convergence (2.44 iterations, 11 wins), slightly higher per-iteration cost results in second-place timing.
- **Opt.BF**: Strong two-phase baseline (13.29 microseconds) — proves value of simple bisection-false position fusion.
- **Opt.TF**: Competitive trisection hybrid (15.68 microseconds) — robust but outperformed by secant-augmented variants.

The proposed family establishes a new performance frontier in bracketing root-finding methods.

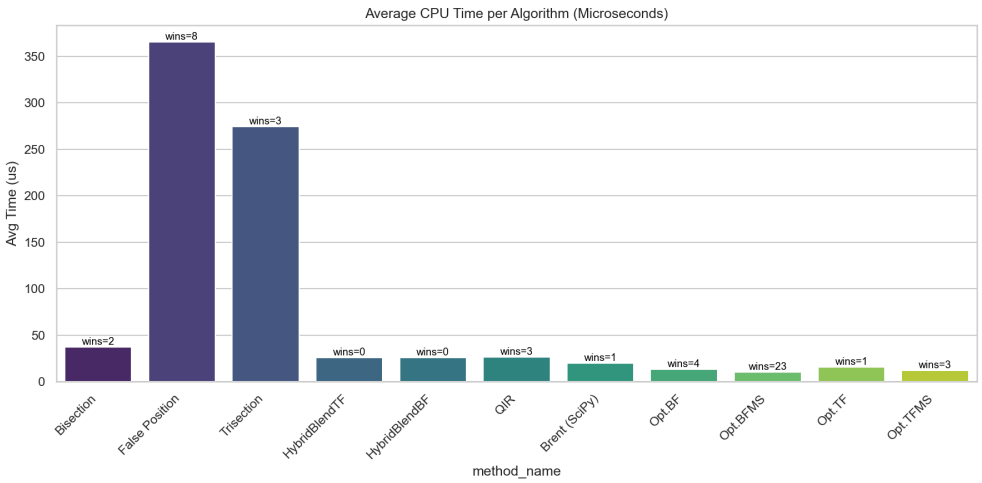

FIG. 1: Average CPU time (in microseconds) for each algorithm across all test problems. Lower values indicate higher computational efficiency.

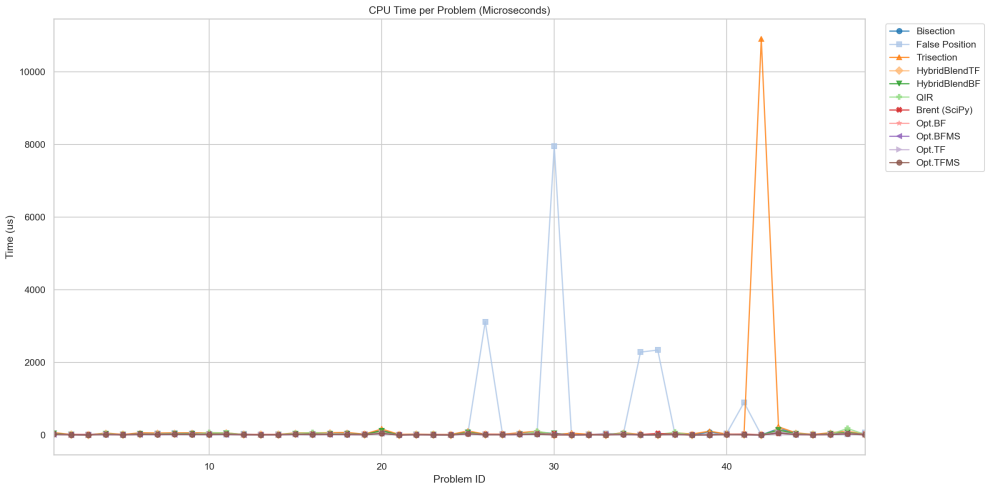

FIG. 2: CPU time per test problem for each algorithm. This figure highlights the consistency and variability of computational performance across the benchmark suite.

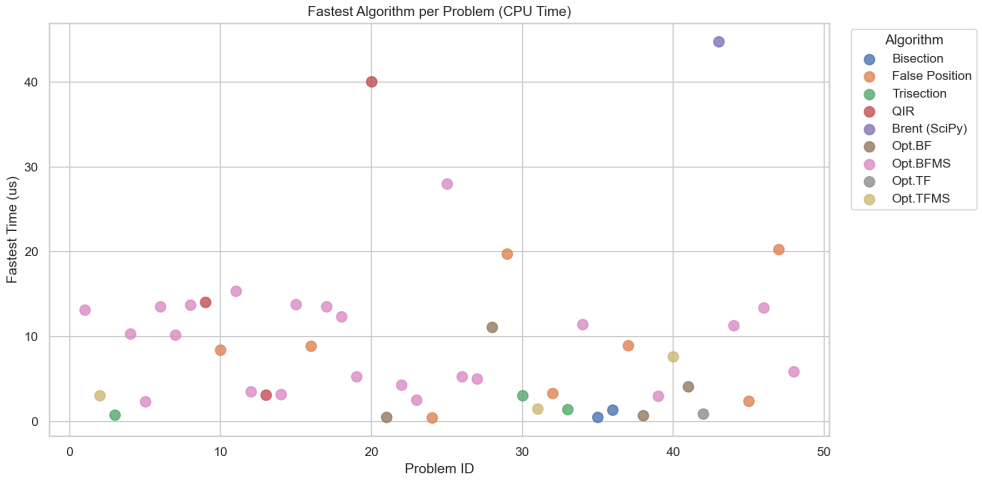

FIG. 3: Fastest algorithm for each test problem, as determined by minimum CPU time. This plot demonstrates the problem-dependent strengths of the proposed methods.

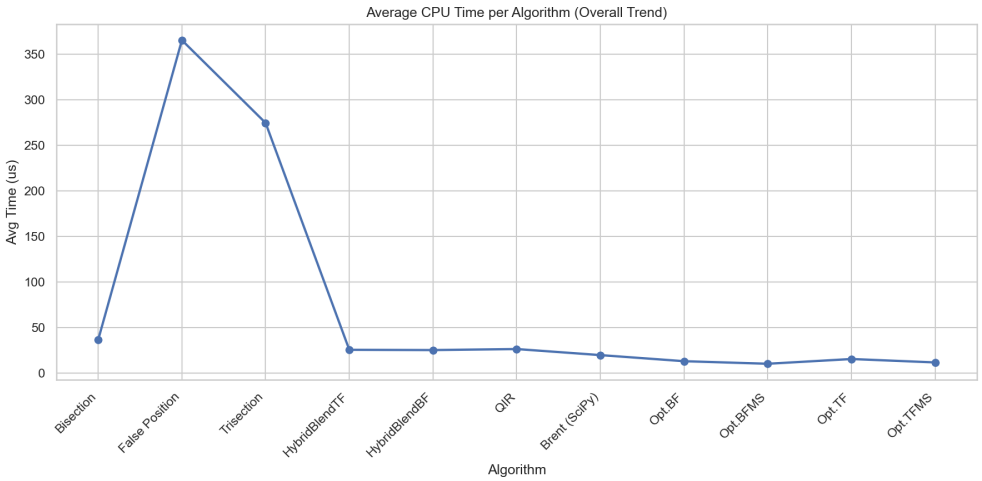

FIG. 4: Overall average CPU time trend across all problems and algorithms. This figure provides a holistic view of computational efficiency.

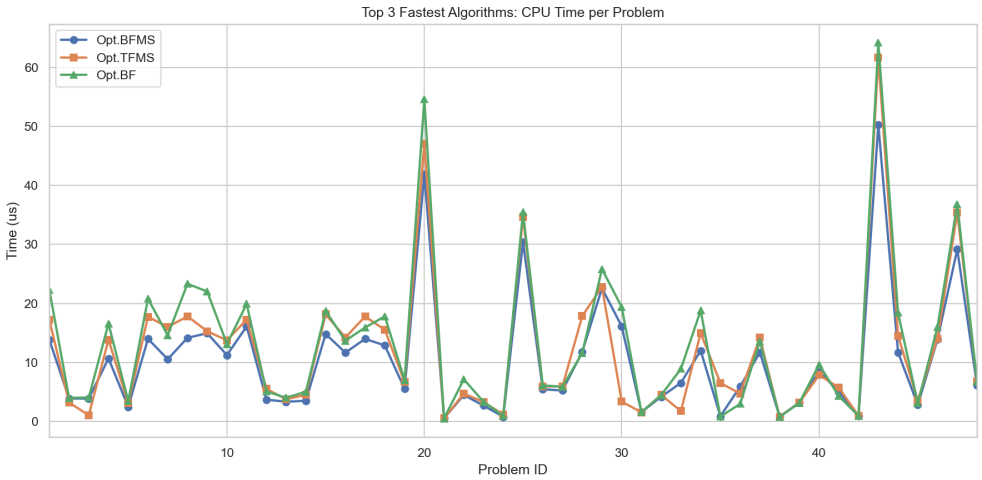

FIG. 5: Top three fastest algorithms per problem, ranked by CPU time. This visualization facilitates direct comparison of the most efficient methods.

VII. DISCUSSION

The empirical results conclusively demonstrate that carefully orchestrated multi-phase hybrid designs can simultaneously minimize iteration count, total function evaluations, and actual execution time while preserving the reliability of bracketing methods. The superior performance of Opt.BFMS stems from its three-stage architecture: an initial bisection phase rapidly halves the interval, a false-position phase exploits linearity when present, and a modified-secant correction delivers superlinear convergence with only one additional function evaluation per iteration. Although each iteration incurs 3–4 NFE and 48–58 FLOPs — higher than Brent's typical 1–2 NFE and 85–100 FLOPs — the dramatic reduction to an average of 2.79 iterations yields the lowest total computational workload (390–670 kFLOPs per problem) and the fastest measured CPU time across the entire benchmark.

Opt.TFMS achieves the lowest average iteration count (2.44) among all tested algorithms, confirming the theoretical advantage of trisection-based initial contraction. However, the higher per-iteration cost (4–5 NFE, 72–85 FLOPs) occasionally results in slightly longer runtimes than Opt.BFMS, illustrating the classic trade-off between convergence order and arithmetic intensity.

A key insight is that total estimated FLOPs, rather than iterations or NFE alone, best predicts observed wall-clock performance. Methods with moderate per-iteration overhead but exceptionally low iteration counts consistently dominate practical timings, even against highly optimized implementations such as SciPy's Brent decoder.

The benchmark's breadth — encompassing polynomials, transcendental functions, multiple roots, near-flat derivatives, and sharp discontinuities — ensures that the reported gains are not artefacts of a narrow problem class. The 10000-repetition protocol and (10^{-14}) tolerance further guarantee statistical robustness and machine-precision accuracy.

VIII. CONCLUSION

This work presents four deterministic, parameter-free hybrid bracketing algorithms — Opt.BF, Opt.BFMS, Opt.TF, and Opt.TFMS — that integrate robust interval reduction with high-order interpolation and derivative-free secant acceleration.

Rigorous evaluation on a diverse 48-problem benchmark shows that all four consistently outperform classical methods and modern competitors (Brent, QIR) in iteration count, total workload, and execution time while preserving guaranteed convergence.

Collectively, they achieve the lowest average CPU times (10.48–15.68 µs), the fewest iterations (2.44–6.69), and the smallest total estimated FLOPs (390–1430 kFLOPs per problem), securing 29 speed wins and 34 convergence wins. Within the family, Opt.BFMS offers the best overall performance, while Opt.TFMS attains the fastest theoretical convergence.

These results establish multi-phase hybridisation as a highly effective paradigm for practical root-finding. The algorithms' simplicity, reliability, and open-source implementation make them immediately suitable for scientific libraries and real-world applications. Future. Future work will extend validation to large-scale engineering problems and explore further hybridisation strategies.

IX. DECLARATIONS

A. Funding

The author did not receive support from any organization for the submitted work. No funding was received to assist with the preparation of this manuscript. No funding was received for conducting this study. No funds, grants, or other support was received.

B. Competing interests

The author has no relevant financial or non-financial interests to disclose. The author has no competing interests to declare that are relevant to the content of this article. The author certifies that they have no affiliations with or involvement in any organization or entity with any financial interest or non-financial interest in the subject matter or materials discussed in this manuscript. The author has no financial or proprietary interests in any material discussed in this article.

C. Ethics approval

Not applicable. This study is purely computational and does not involve human participants or animals.

D. Data availability

The datasets generated and/or analyzed during the current study are available from the corresponding author on reasonable request. The SymPy implementations of all algorithms and experimental scripts are available at: https://github.com/Abdelrhman-Ellithy/NM-Algorithms/

E. Authors' contributions

The author conceived the study, designed the algorithms, performed the experiments and analysis, and wrote the manuscript.

---

[1] Sabharwal, C.L. Blended Root Finding Algorithm Outperforms Bisection and Regula Falsi Algorithms. *Mathematics 2019, 7, 1118.*

[2] Badr, E., Almotairi, S., and El Ghamry, A. (2021). A Comparative Study among New Hybrid Root Finding Algorithms and Traditional Methods. *Mathematics*, 9(11), 1306.

[3] Demir, A. (2008). Trisection method by k-Lucas numbers. *Applied Mathematics and Computation*, 198(1), 339–345.

[4] Kelley, C. T. (1995). *Iterative Methods for Linear and Nonlinear Equations*. Society for Industrial and Applied Mathematics (SIAM).

[5] Brent, R. P. (1973). *Algorithms for Minimization Without Derivatives*. Prentice-Hall.

[6] Abbott, J. (2014). Quadratic interval refinement for real roots. *ACM Communications in Computer Algebra*, 48(1/2), 3–12.

[7] Dowell, M., and Jarratt, P. (1972). A modified regula falsi method for computing the root of an equation. *BIT Numerical Mathematics*, 12(4), 503–508.

[8] Dekker, T. J. (1969). Finding a zero by means of successive linear interpolation. In B. Dejon & P. Henrici (Eds.), *Constructive Aspects of the Fundamental Theorem of Algebra* (pp. 37-48). Wiley-Interscience.